

Predicting_Customer_Churn_in_Telecom_Operators

March 30, 2020

1 Predicting Customer Churn in Telecommunication Operators

Customer turnover refers to a decision made by the customer on the term of business relationship. Customer loyalty and customer turnover always add up to 100%. If a company has a 60% loyalty rate, then customer loss taxes are 40%. According to the 80/20 customer profitability rule, 20% of customers are generating 80% of revenue. Therefore, it is very important to predict the users who are likely to abandon the business relationship and the factors that affect how the customer's decisions. In this project, I predicted Customer Churn at a Telecommunications Operator using python and frameworks (Pandas, Numpy, scipy and Scikit-Learn).

```
[8]: # Importing libraries and frameworks
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

import scipy.stats

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler

from sklearn.feature_selection import f_classif, chi2
from sklearn.feature_selection import SelectKBest

import random

import folium
import os
from IPython.display import display
from IPython.display import Image

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import confusion_matrix, classification_report, \
    accuracy_score
```

```
from sklearn.externals import joblib

import warnings
warnings.filterwarnings("ignore")
```

1.1 Importing dataset

```
[2]: # Importing train dataset
df_train = pd.read_csv("data/projeto4_telecom_treino.csv")
df_test = pd.read_csv("data/projeto4_telecom_teste.csv")
```

1.2 Exploratory Analysis

```
[82]: # Checking dataset features
print(df_train.shape[1])
df_train.head()
```

21

```
[82]:  id state  account_length      area_code international_plan voice_mail_plan \
0    1    KS           128  area_code_415             no             yes
1    2    OH           107  area_code_415             no             yes
2    3    NJ           137  area_code_415             no             no
3    4    OH            84  area_code_408             yes             no
4    5    OK            75  area_code_415             yes             no
```

```
      number_vmail_messages  total_day_minutes  total_day_calls  \
0                        25             265.1             110
1                        26             161.6             123
2                        0             243.4             114
3                        0             299.4              71
4                        0             166.7             113
```

```
      total_day_charge  ...  total_eve_calls  total_eve_charge  \
0             45.07  ...             99             16.78
1             27.47  ...             103             16.62
2             41.38  ...             110             10.30
3             50.90  ...              88              5.26
4             28.34  ...             122             12.61
```

```
      total_night_minutes  total_night_calls  total_night_charge  \
0             244.7             91             11.01
1             254.4             103             11.45
2             162.6             104              7.32
3             196.9             89              8.86
```

4	186.9	121	8.41
---	-------	-----	------

	total_intl_minutes	total_intl_calls	total_intl_charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73

	number_customer_service_calls	churn
0	1	no
1	1	no
2	0	no
3	2	no
4	3	no

[5 rows x 21 columns]

```
[83]: # Checking for missing values
print(pd.isna(df_train).any().any())
print(pd.isna(df_test).any().any())
print(pd.isnull(df_train).any().any())
print(pd.isnull(df_test).any().any())
```

False
False
False
False

```
[84]: # datasets size
print(df_train.shape)
print(df_test.shape)
```

(3333, 21)
(1667, 21)

```
[207]: df_train.columns
```

```
[207]: Index(['id', 'state', 'account_length', 'area_code', 'international_plan',
          'voice_mail_plan', 'number_vmail_messages', 'total_day_minutes',
          'total_day_calls', 'total_day_charge', 'total_eve_minutes',
          'total_eve_calls', 'total_eve_charge', 'total_night_minutes',
          'total_night_calls', 'total_night_charge', 'total_intl_minutes',
          'total_intl_calls', 'total_intl_charge',
          'number_customer_service_calls', 'churn'],
          dtype='object')
```

```
[85]: # Compute numerical data summary statistics
df_train_num = df_train[['account_length', 'number_vmail_messages',
    → 'total_day_minutes',
    'total_day_calls', 'total_day_charge',
    → 'total_eve_minutes',
    'total_eve_calls', 'total_eve_charge',
    → 'total_night_minutes',
    'total_night_calls', 'total_night_charge',
    → 'total_intl_minutes',
    'total_intl_calls', 'total_intl_charge',
    → 'number_customer_service_calls']].copy()

df_describe = pd.concat([df_train_num.describe().T,
    df_train_num.mad().rename('mad'),
    df_train_num.skew().rename('skew'),
    df_train_num.kurt().rename('kurt'),
    df_train_num.median().rename('median')
    ], axis=1).T

display(df_describe)
```

	account_length	number_vmail_messages	total_day_minutes	\
count	3333.000000	3333.000000	3333.000000	
mean	101.064806	8.099010	179.775098	
std	39.822106	13.688365	54.467389	
min	1.000000	0.000000	0.000000	
25%	74.000000	0.000000	143.700000	
50%	101.000000	0.000000	179.400000	
75%	127.000000	20.000000	216.400000	
max	243.000000	51.000000	350.800000	
mad	31.821440	11.719778	43.523455	
skew	0.096606	1.264824	-0.029077	
kurt	-0.107836	-0.051129	-0.019940	
median	101.000000	0.000000	179.400000	

	total_day_calls	total_day_charge	total_eve_minutes	total_eve_calls	\
count	3333.000000	3333.000000	3333.000000	3333.000000	
mean	100.435644	30.562307	200.980348	100.114311	
std	20.069084	9.259435	50.713844	19.922625	
min	0.000000	0.000000	0.000000	0.000000	
25%	87.000000	24.430000	166.600000	87.000000	
50%	101.000000	30.500000	201.400000	100.000000	
75%	114.000000	36.790000	235.300000	114.000000	

max	165.000000	59.640000	363.700000	170.000000
mad	15.944943	7.398914	40.469244	15.860332
skew	-0.111787	-0.029083	-0.023877	-0.055563
kurt	0.243182	-0.019812	0.025630	0.206156
median	101.000000	30.500000	201.400000	100.000000

	total_eve_charge	total_night_minutes	total_night_calls	\
count	3333.000000	3333.000000	3333.000000	
mean	17.083540	200.872037	100.107711	
std	4.310668	50.573847	19.568609	
min	0.000000	23.200000	33.000000	
25%	14.160000	167.000000	87.000000	
50%	17.120000	201.200000	100.000000	
75%	20.000000	235.300000	113.000000	
max	30.910000	395.000000	175.000000	
mad	3.439937	40.410387	15.690341	
skew	-0.023858	0.008921	0.032500	
kurt	0.025487	0.085816	-0.072020	
median	17.120000	201.200000	100.000000	

	total_night_charge	total_intl_minutes	total_intl_calls	\
count	3333.000000	3333.000000	3333.000000	
mean	9.039325	10.237294	4.479448	
std	2.275873	2.791840	2.461214	
min	1.040000	0.000000	0.000000	
25%	7.520000	8.500000	3.000000	
50%	9.050000	10.300000	4.000000	
75%	10.590000	12.100000	6.000000	
max	17.770000	20.000000	20.000000	
mad	1.818555	2.184712	1.881093	
skew	0.008886	-0.245136	1.321478	
kurt	0.085663	0.609185	3.083589	
median	9.050000	10.300000	4.000000	

	total_intl_charge	number_customer_service_calls
count	3333.000000	3333.000000
mean	2.764581	1.562856
std	0.753773	1.315491
min	0.000000	0.000000
25%	2.300000	1.000000
50%	2.780000	1.000000
75%	3.270000	2.000000
max	5.400000	9.000000
mad	0.589880	1.052532
skew	-0.245287	1.091359
kurt	0.609610	1.730914
median	2.780000	1.000000

```
[86]: # Compute categorical data summary statistics
df_train_cat = df_train[['state', 'area_code', 'international_plan', 'voice_mail_plan', 'churn']].copy()
df_train_cat.describe()
```

```
[86]:      state      area_code international_plan voice_mail_plan churn
count  3333      3333      3333      3333      3333
unique   51         3         2         2         2
top      WV  area_code_415      no      no      no
freq    106      1655      3010      2411      2850
```

1.2.1 Univariate analysis


Checking churn variable distribution and proportion

```
[214]: # churn values and proportion
df = pd.DataFrame(pd.Categorical(df_train.churn).describe())
display(df)

# churn variable barplot
sns.set(style="whitegrid")
sns.barplot(x=['no', 'yes'], y=df_train.churn.value_counts().values).
    set_title('Churn barplot')
```

	counts	freqs
categories		
no	2850	0.855086
yes	483	0.144914

```
[214]: Text(0.5, 1.0, 'Churn barplot')
```



As noticed on table and graph above, churn feature is unbalanced. About 14% of customers stopped using the telecom service and 85% still using it.

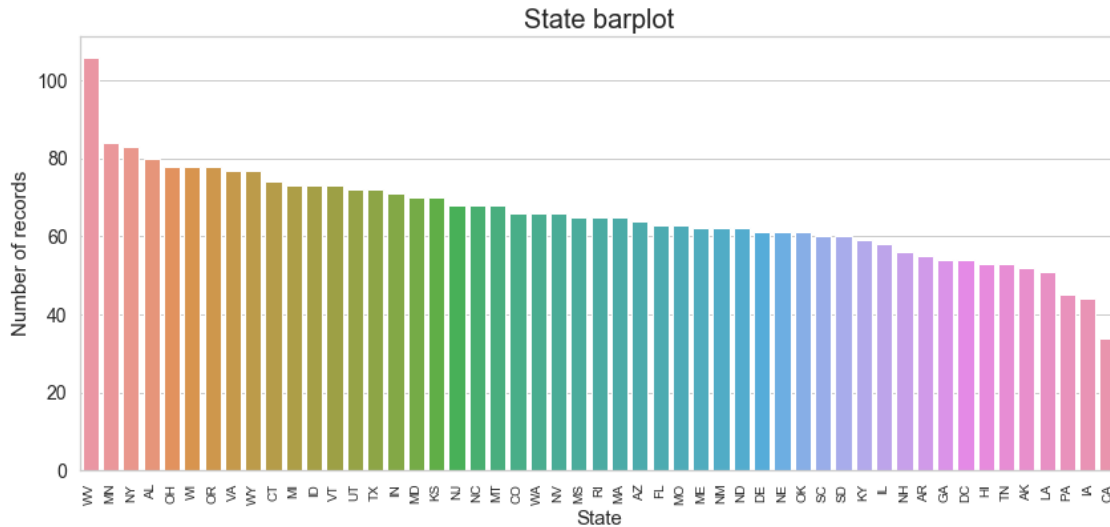
Categorical variables

State

```
[232]: # State variable barplot
plt.figure(figsize=(14,6))

sns.barplot(x=df_train.state.value_counts().index,
            y=df_train.state.value_counts().values)

plt.title('State barplot', size=20)
plt.xlabel('State', size=15)
plt.ylabel('Number of records', size=15)
plt.yticks(fontsize=14)
plt.xticks(fontsize=10, rotation=90)
plt.show()
```



The state with the highest frequency is the West Virginia and the lowest frequency is the California. Other states with a large number of records are New York and Minnesota. ##### Number of records by State shown on the map below:

```
[9]: # Map graph
# Load the shape of the zone (US states)
state_geo = os.path.join('', 'us-states.json')

# state data
state_data = pd.DataFrame({'state': df_train.state.value_counts().index,
                           'count': df_train.state.value_counts().values})

# Initialize the map:
m = folium.Map(location=[37, -102], zoom_start=5)

# Add the color for the choropleth:
m.choropleth(
    geo_data=state_geo,
    name='Number of records by state',
    data=state_data,
    columns=['state', 'count'],
    key_on='feature.id',
    fill_color='YlGn',
    fill_opacity=0.7,
    line_opacity=0.2,
    legend_name='Number of records'
)
folium.LayerControl().add_to(m)

# Save to html
```



```

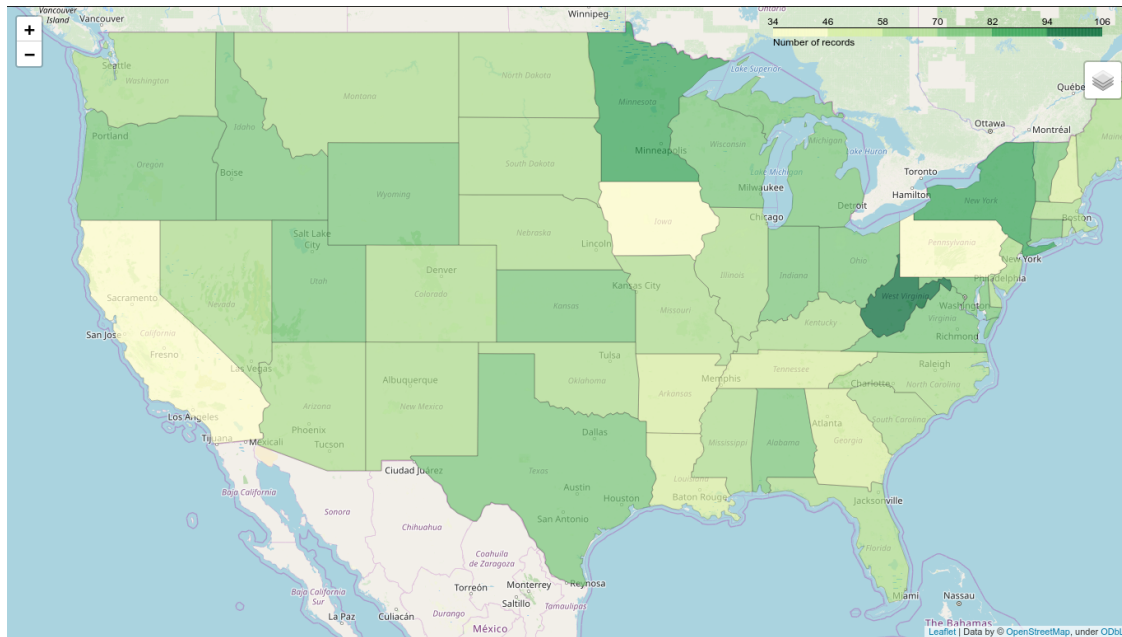
m.save('#registers_by_map.html')

display(m)

# Loading map image
# Image(filename='records_by_state.png')

```

[9]:



Area code

```

[27]: # area_code values and proportion
plt.figure(figsize=(10,6))


df = pd.DataFrame(pd.Categorical(df_train.area_code).describe())
display(df)

# area_code variable barplot
sns.set(style="whitegrid")
sns.barplot(x=df_train.area_code.value_counts().index,
            y=df_train.area_code.value_counts().values,
            palette='Set1')

plt.title('Area code barplot', size=20)
plt.xlabel('Area code', size=15)
plt.ylabel('Number of records', size=15)
plt.yticks(fontsize=14)
plt.xticks(fontsize=14)
plt.show()

```

	counts	freqs
categories		
area_code_408	838	0.251425
area_code_415	1655	0.496550
area_code_510	840	0.252025



Code area 415 has the largest number of records. ##### International plan

```
[29]: # international_plan values and proportion
plt.figure(figsize=(8,4))

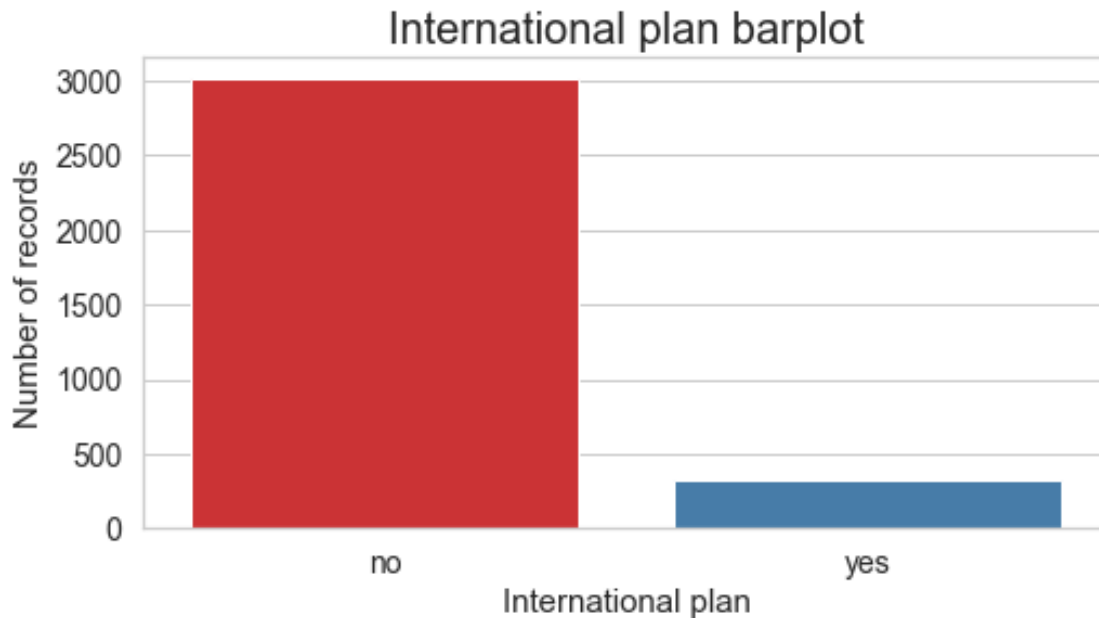
df = pd.DataFrame(pd.Categorical(df_train.international_plan).describe())
display(df)

# international_plan variable barplot
sns.set(style="whitegrid")
sns.barplot(x=df_train.international_plan.value_counts().index,
            y=df_train.international_plan.value_counts().values,
            palette='Set1')

plt.title('International plan barplot', size=20)
plt.xlabel('International plan', size=15)
plt.ylabel('Number of records', size=15)
plt.yticks(fontsize=14)
plt.xticks(fontsize=14)
```

```
plt.show()
```

	counts	freqs
no	3010	0.90309
yes	323	0.09691



Most customers do not have international plan. ##### Voice mail plan

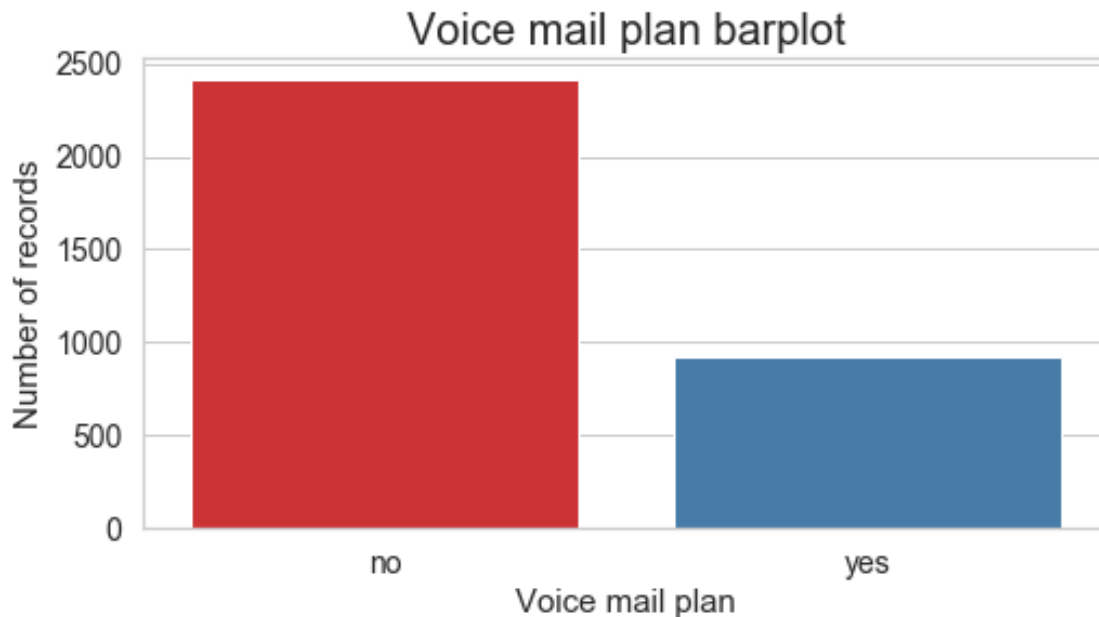
```
[30]: # voice_mail_plan values and proportion
plt.figure(figsize=(8,4))

df = pd.DataFrame(pd.Categorical(df_train.voice_mail_plan).describe())
display(df)

# voice_mail_plan variable barplot
sns.set(style="whitegrid")
sns.barplot(x=df_train.voice_mail_plan.value_counts().index,
            y=df_train.voice_mail_plan.value_counts().values,
            palette='Set1')

plt.title('Voice mail plan barplot', size=20)
plt.xlabel('Voice mail plan', size=15)
plt.ylabel('Number of records', size=15)
plt.yticks(fontsize=14)
plt.xticks(fontsize=14)
plt.show()
```

	counts	freqs
categories		
no	2411	0.723372
yes	922	0.276628



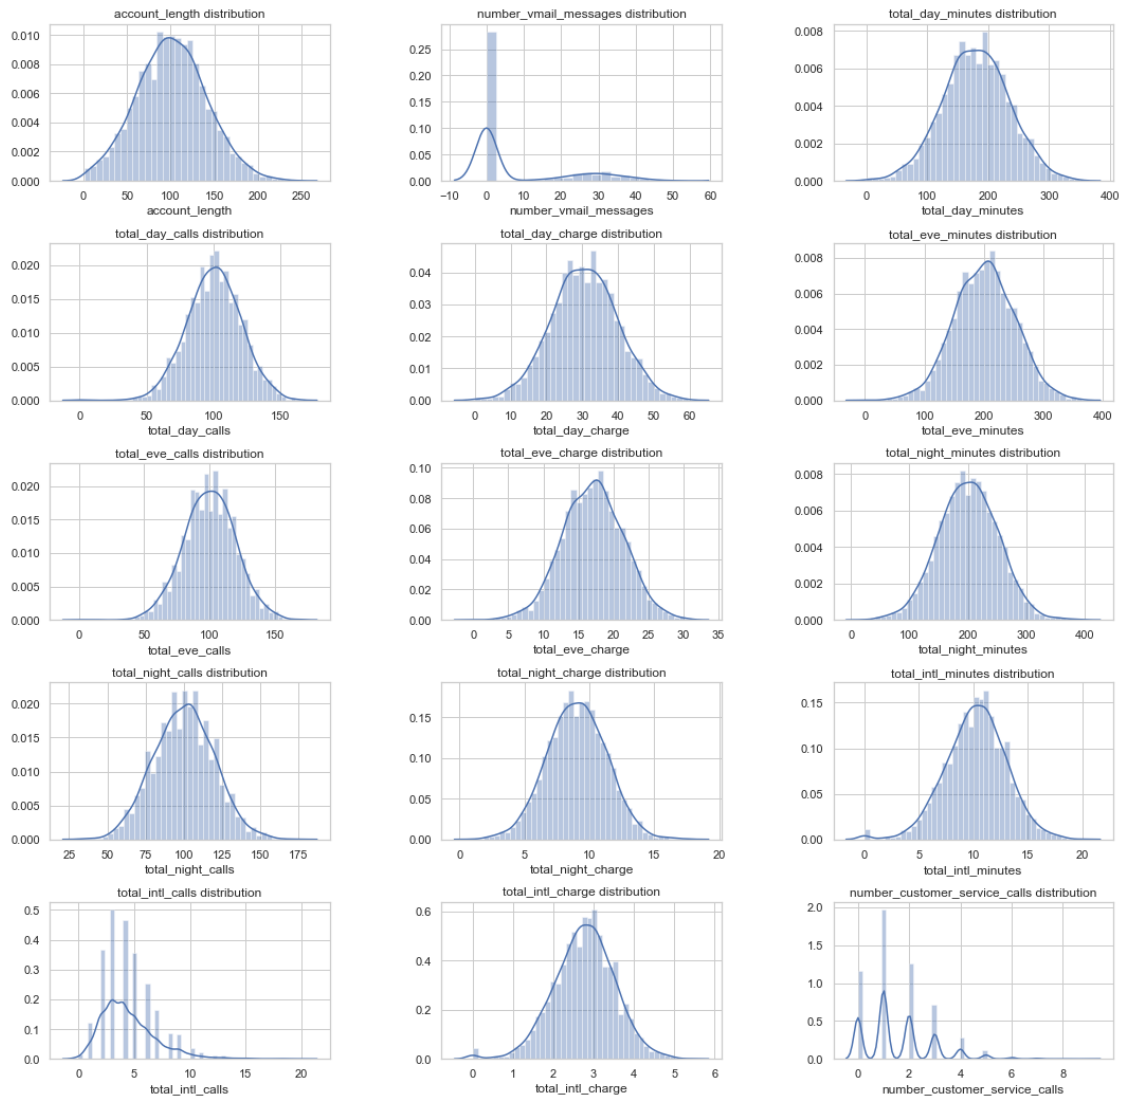
Most customers do not have voice mail plan. ##### Numerical variables ##### Checking numerical features distributions

```
[40]: # Features histograms and kde
df = df_train_num.copy()
fig, axs = plt.subplots(ncols=3, nrows=5)
plt.subplots_adjust(hspace = 0.4, wspace = 0.4)
fig.set_size_inches(18, 18, forward=True)

count = 0

for i in range(5):
    for j in range(3):
        plt.sca(ax=axs[i][j])
        if count < df.shape[1]:
            col = df.columns[count]
            sns.distplot(df[col]).set_title(col + ' distribution')
        else:
            break

    count += 1
```



“account_lenght”, “ total_day_minutes”, “ total_day_calls”, “ total_day_charge”, “ total_eve_minutes”, “ total_eve_calls”, “ total_eve_charge”, “ total_night_minutes”, “ total_night_calls”, “ total_night_charge”, “ total_intl_minutes”, “ total_intl_charge” seem to have a normal distribution. “number_vmail_messages” has a bimodal distribution; “total_intl_calls” has a exponential distribution and “number_customer_service_calls” has a multimodal distribution.

```
[42]: # Features boxplot
df = df_train_num.copy()
fig, axs = plt.subplots(ncols=3, nrows=5)
plt.subplots_adjust(hspace = 0.4, wspace = 0.4)
fig.set_size_inches(18, 18, forward=True)

count = 0

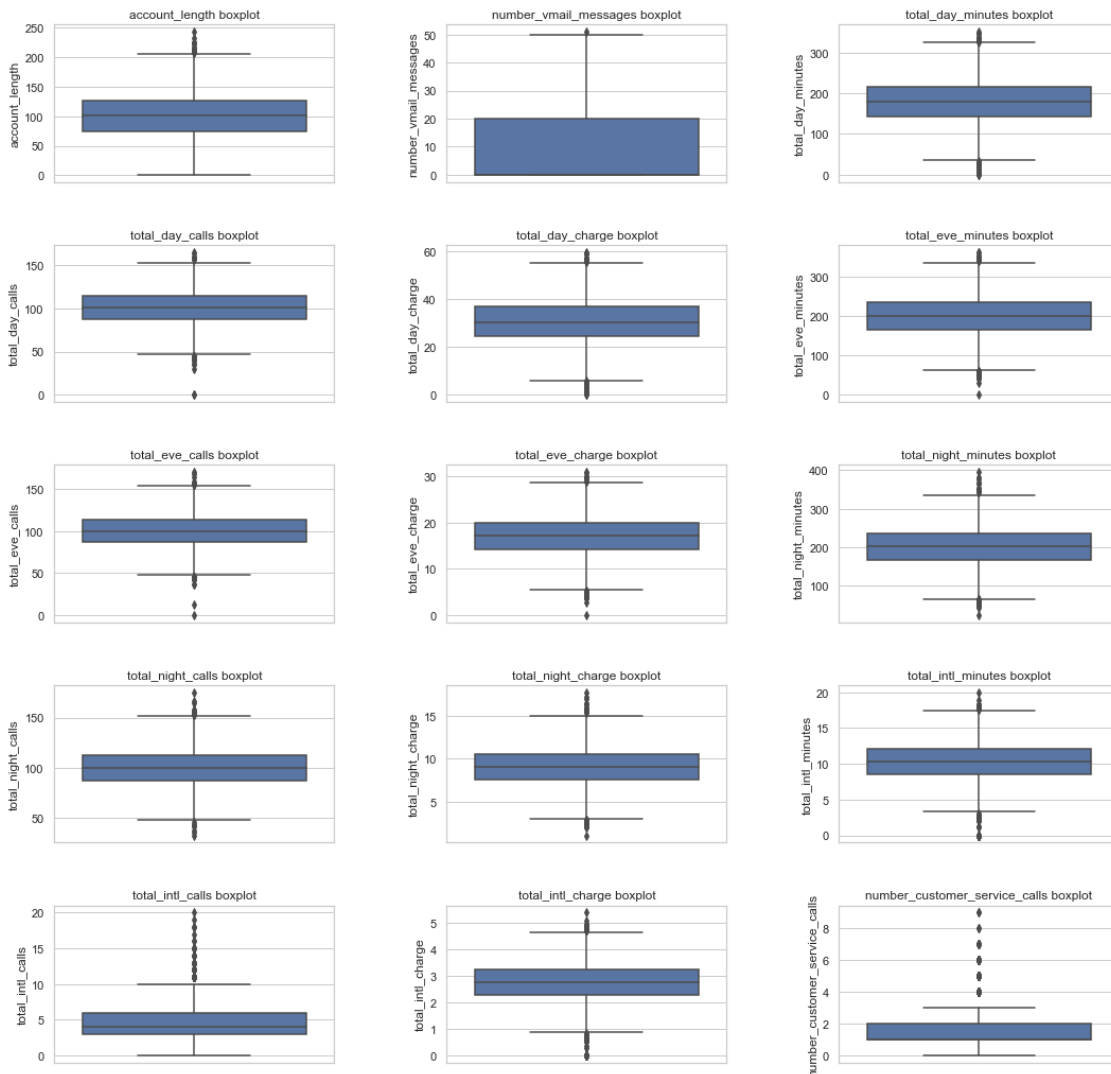
for i in range(5):
```

```

for j in range(3):
    plt.sca(ax=axes[i][j])
    if count < df.shape[1]:
        col = df.columns[count]
        sns.boxplot(y=df[col]).set_title(col + ' boxplot')
    else:
        break

count +=1

```



“total_intl_calls” and “number_customer_service” have a large number of outliers. ### Bivariate analysis ##### Checking correlation between numerical variables

```

[16]: # heat map of correlation values
df = df_train_num.copy()

```

```

df['churn'] = df_train.churn.apply(lambda x: 0 if x=='no' else 1)


corr = df.corr()
fig, ax = plt.subplots(figsize=(16,16))

g = sns.heatmap(corr, annot=True, ax=ax, square=True, linewidth=0.5)

plt.yticks(rotation=0)
g.set_xticklabels(g.get_xticklabels(), rotation=90)
ax.set_ylim([len(corr) + 0.5, 0])
ax.set_xlim([-0.5, len(corr)])

```

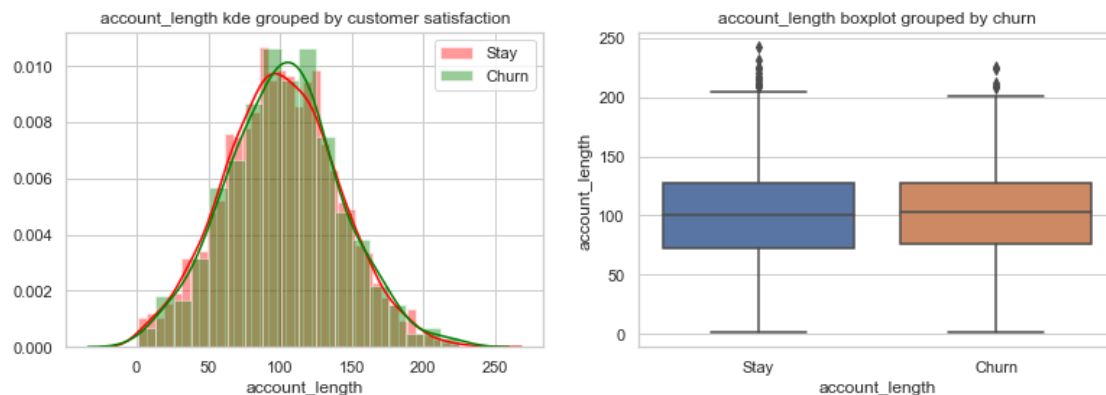
[16]: (-0.5, 16)

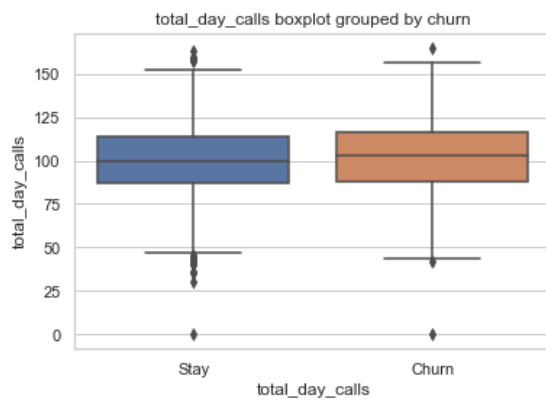

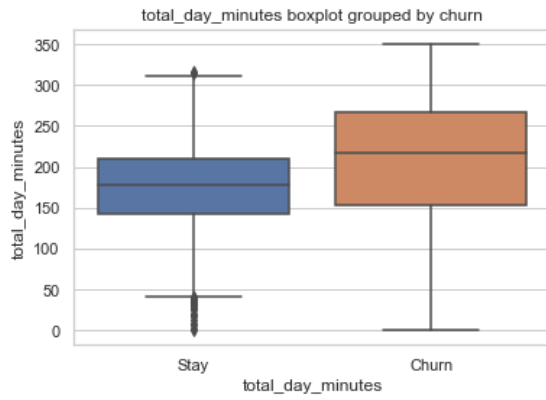
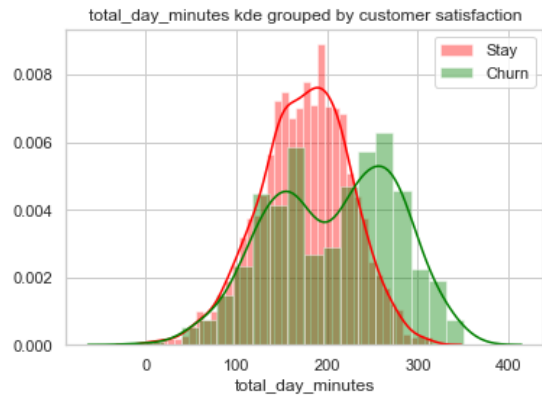
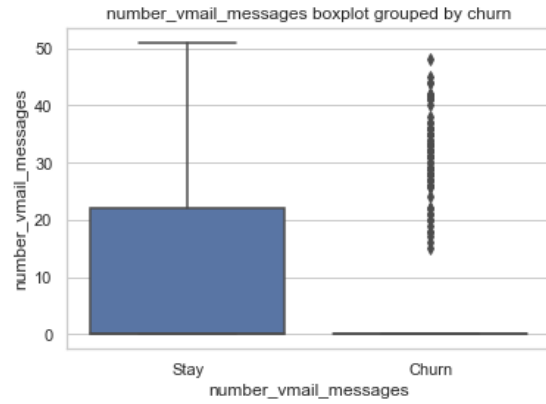



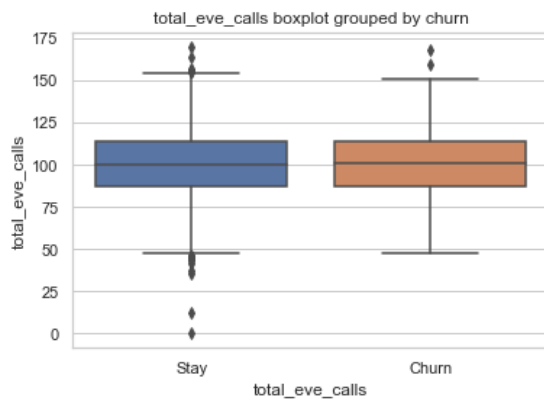


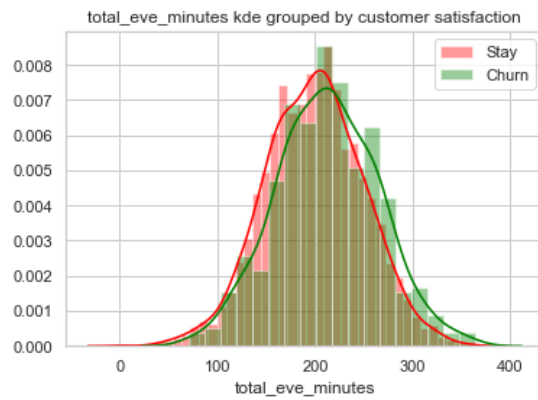


There is a high correlation between the following pairs of variables: “total_day_minutes” and “total_day_charge”, “total_night_minutes” and “total_night_charge”, “total_eve_minutes” and “total_eve_charge”, “total_intl_minutes” and “total_intl_charge”. ##### Getting numerical features distribution grouped by target variable (churn)

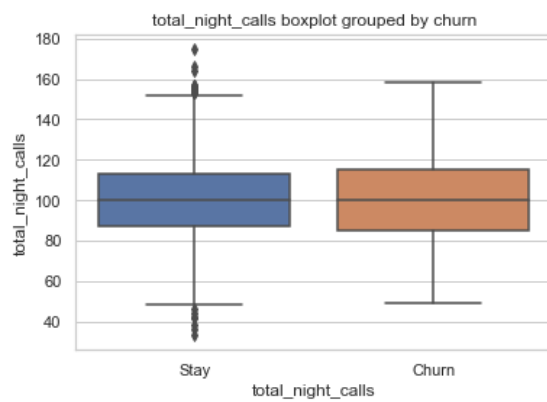
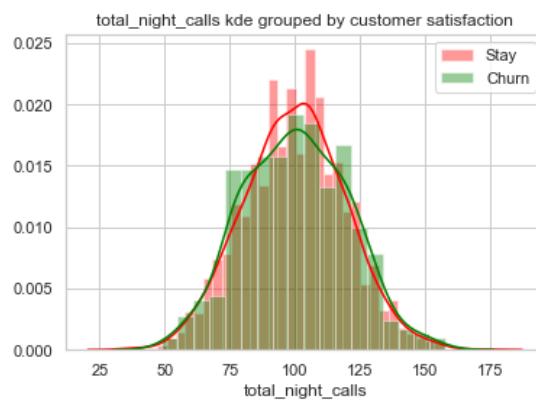
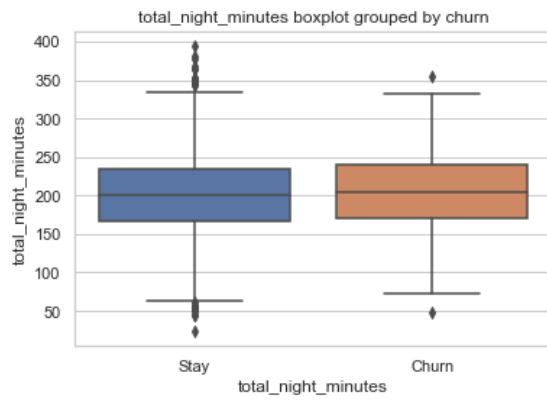
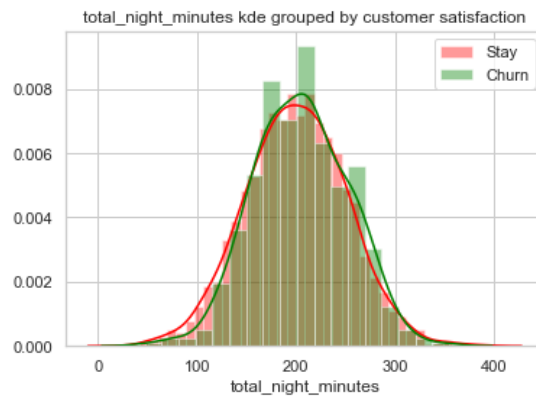


```
[50]: # Getting numerical features distribution grouped by churn variable
df = df_train_num
df['churn'] = df_train.churn.copy()

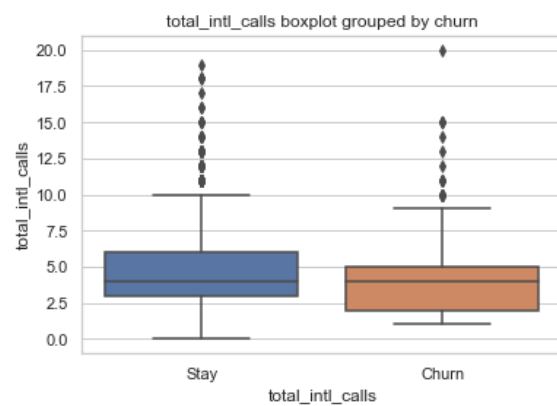



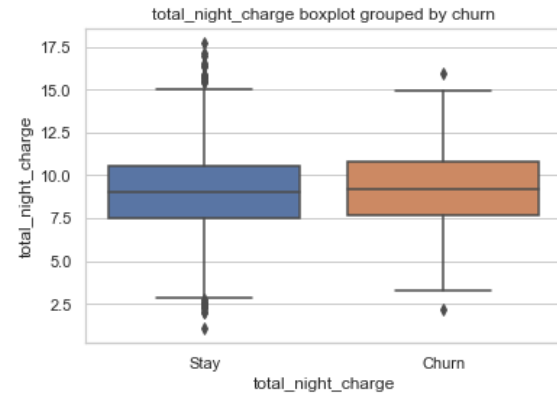

for col in df.drop("churn", axis=1).columns:
    fig, axs = plt.subplots(ncols=2)
    fig.set_size_inches(13, 4, forward=True)
    sns.distplot(df[df.churn == 'no'][col], color='red', label='Stay',
    →ax=axs[0], bins = 40)
    sns.distplot(df[df.churn == 'yes'][col], color='green', label='Churn',
    →ax=axs[0], bins = 18)
    axs[0].legend()
    axs[0].set_xlabel(col)
    axs[0].set_title(col + ' kde grouped by customer satisfaction')
    sns.boxplot(y=col, x="churn", data = df, ax=axs[1])
    axs[1].set_xlabel(col)
    axs[1].set_xticklabels(['Stay', 'Churn'])
    axs[1].set_title(col + ' boxplot grouped by churn')
plt.show()
```

















Getting categorical features barplot grouped by churn variable

```
[87]: # categorical variables boxplot grouped by churn
df = df_train_cat.drop('state', axis=1)
df['churn'] = df_train.churn.copy()

fig, axs = plt.subplots(ncols=3)
plt.subplots_adjust(hspace = 0.4, wspace = 0.4)
fig.set_size_inches(18, 4, forward=True)


count = 0

for j in range(3):
    plt.sca(ax=axs[j])
    if count < df.shape[1]:
        col = df.columns[count]
        df_g = pd.DataFrame(df.groupby([col, 'churn']).size()).reset_index()
        sns.barplot(x=col, y=0, hue="churn", data=df_g, palette="rocket")
```

```

    axs[j].set_xlabel(col)
    axs[j].set_title(col + ' barplot grouped by churn')
else:
    break
count +=1

```



Most of the customers who remained do not have an international plan. ##### Churn by state

```

[11]: # Map graph exibing churn proportion
# Load the shape of the zone (US states)
state_geo = os.path.join('', 'us-states.json')

# state data
state_data = pd.DataFrame({'state': df_train[df_train.churn == 'yes'].state.
    →value_counts().index,
                           'churn_prop': df_train[df_train.churn == 'yes'].
    →state.value_counts().values/ \
                           df_train.state.value_counts().values})

state_data

# Initialize the map:
m = folium.Map(location=[37, -102], zoom_start=5)

# Add the color for the choropleth:
m.choropleth(
    geo_data=state_geo,
    name='Churn proportion by state',
    data=state_data,
    columns=['state', 'churn_prop'],
    key_on='feature.id',
    fill_color='PuRd',
    fill_opacity=0.7,
    line_opacity=0.2,
    legend_name='Churn proportion'
)
folium.LayerControl().add_to(m)

```

```

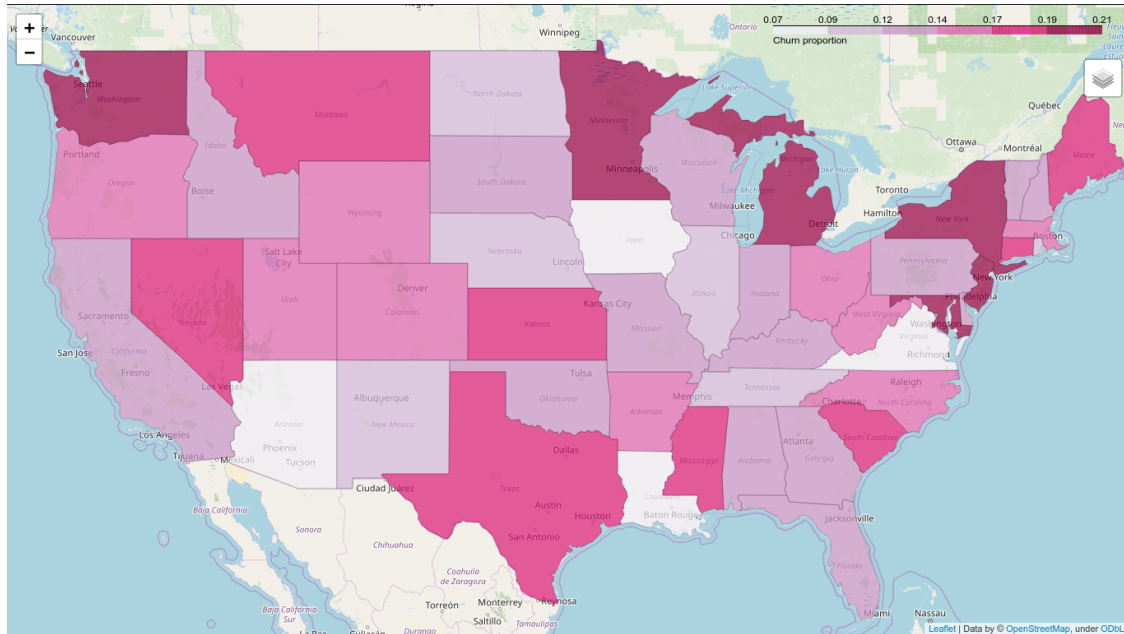
# Save to html
m.save('#churn_proportion.html')

display(m)

# # Loading map image
# Image(filename='churn_proportion_by_state.png')

```

[11]:



Checking categorical features

```

[88]: df_new_train_cat = df_train_cat.copy()

# Converting to binary
df_new_train_cat.international_plan = df_new_train_cat.international_plan.
    →apply(lambda x: 0 if x=='no' else 1)
df_new_train_cat.voice_mail_plan = df_new_train_cat.voice_mail_plan.
    →apply(lambda x: 0 if x=='no' else 1)
df_new_train_cat.churn = df_new_train_cat.churn.apply(lambda x: 0 if x=='no' else 1)

dic = {}
count = 1

# Changing state variable
for state in df_new_train_cat.state.unique():
    dic[state] = count

```

```

count += 1

df_new_train_cat.state = df_new_train_cat.state.apply(lambda x: dic[x])

# Changing area_code variable
df_new_train_cat.area_code[df_new_train_cat.area_code == 'area_code_408'] = 1
df_new_train_cat.area_code[df_new_train_cat.area_code == 'area_code_510'] = 2
df_new_train_cat.area_code[df_new_train_cat.area_code == 'area_code_415'] = 3

```

```

[89]: # heat map of correlation values
df = df_new_train_cat.copy()

corr = df.corr()
fig, ax = plt.subplots(figsize=(8,8))

g = sns.heatmap(corr, annot=True, ax=ax, square=True, linewidth=0.5)


plt.yticks(rotation=0)
g.set_xticklabels(g.get_xticklabels(), rotation=90)
ax.set_ylim([len(corr) + 0.5, 0])
ax.set_xlim([-0.5, len(corr)])

```

```

[89]: (-0.5, 4)

```

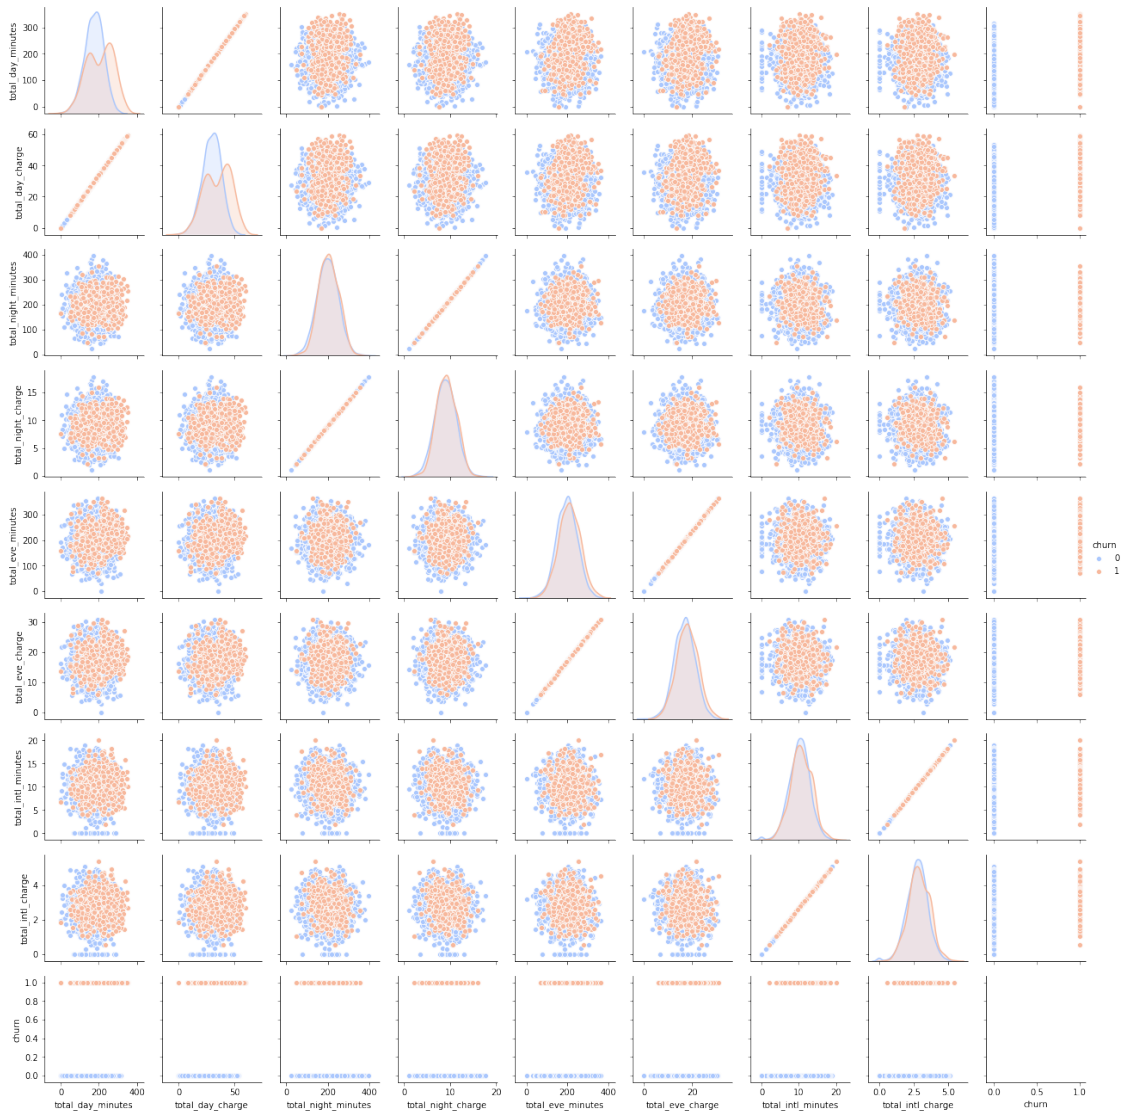



1.2.2 Multivariate analysis

```
[17]: # pairplot matrix grouped by churn variable
df = df_train[["total_day_minutes", "total_day_charge", "total_night_minutes",
               "total_night_charge", "total_eve_minutes",
               → "total_eve_charge",
               "total_intl_minutes", "total_intl_charge", "churn"]].copy()

df['churn'] = df_train.churn.apply(lambda x: 0 if x=='no' else 1)

g = sns.pairplot(df, hue='churn', palette='coolwarm')
g.fig.set_size_inches(18,18)
```



```
[90]: # df_new_train_cat = df_train_cat
df_new_train_cat = df_train_cat.copy()
df_new_train_cat.international_plan = df_new_train_cat.international_plan.
    ↳apply(lambda x: 0 if x=='no' else 1)
df_new_train_cat.voice_mail_plan = df_new_train_cat.voice_mail_plan.
    ↳apply(lambda x: 0 if x=='no' else 1)
df_new_train_cat.churn = df_new_train_cat.churn.apply(lambda x: 0 if x=='no'
    ↳else 1)
```

```
[91]: df_new_train_cat.head()
```

```
[91]: state    area_code  international_plan  voice_mail_plan  churn
0    KS  area_code_415                0                1        0
1    OH  area_code_415                0                1        0
```

2	NJ	area_code_415	0	0	0
3	OH	area_code_408	1	0	0
4	OK	area_code_415	1	0	0

```
[96]: # Checking relation between categorical variables grouped by churn
values = df_train_cat[['state', 'international_plan', 'voice_mail_plan',
    → 'churn']] \
    .groupby(['international_plan', 'voice_mail_plan', 'churn']).count()

dic = {'international_plan': [],
      'voice_mail_plan': [],
      'churn': [],
      'count': []}

for a_tuple in values.index:
    dic['international_plan'].append(a_tuple[0])
    dic['voice_mail_plan'].append(a_tuple[1])
    dic['churn'].append(a_tuple[2])
    dic['count'].append(values.loc[a_tuple].values[0])

display(pd.DataFrame(dic).sort_values(by='count', ascending=False))
```

	international_plan	voice_mail_plan	churn	count
0	0	0	0	1878
2	0	1	0	786
1	0	0	1	302
4	1	0	0	130
5	1	0	1	101
6	1	1	0	56
3	0	1	1	44
7	1	1	1	36

1.3 Feature Engineering

Converting train categorical variables to numerical

```
[92]: # converting test categorical variables to numerical
train_data = df_train.drop('id', axis=1).copy()

# Converting to binary
train_data.international_plan = train_data.international_plan.apply(lambda x: 0 if
    → x=='no' else 1)
train_data.voice_mail_plan = train_data.voice_mail_plan.apply(lambda x: 0 if
    → x=='no' else 1)
train_data.churn = train_data.churn.apply(lambda x: 0 if x=='no' else 1)
```

```

dic = {}
count = 1

# Changing state variable
for state in train_data.state.unique():
    dic[state] = count
    count += 1

train_data.state = train_data.state.apply(lambda x: dic[x])

# Changing area_code variable
train_data.area_code[train_data.area_code == 'area_code_408'] = 1
train_data.area_code[train_data.area_code == 'area_code_510'] = 2
train_data.area_code[train_data.area_code == 'area_code_415'] = 3

# Saving dataset
train_data.to_csv("data/train_data.csv", index=False)

```

Converting test categorical variables to numerical

```

[93]: # converting test categorical variables to numerical
test_data = df_test.drop('id', axis=1).copy()

# Converting to binary
test_data.international_plan = test_data.international_plan.apply(lambda x: 0 if x == 'no' else 1)
test_data.voice_mail_plan = test_data.voice_mail_plan.apply(lambda x: 0 if x == 'no' else 1)
test_data.churn = test_data.churn.apply(lambda x: 0 if x == 'no' else 1)

dic = {}
count = 1

# Changing state variable
for state in test_data.state.unique():
    dic[state] = count
    count += 1

test_data.state = test_data.state.apply(lambda x: dic[x])

# Changing area_code variable
test_data.area_code[test_data.area_code == 'area_code_408'] = 1
test_data.area_code[test_data.area_code == 'area_code_510'] = 2
test_data.area_code[test_data.area_code == 'area_code_415'] = 3

# Saving dataset
test_data.to_csv("data/test_data.csv", index=False)

```

Removing high correlated features

```
[94]: high_corr_features = [col for col in train_data.columns if 'minute' in col]
      high_corr_features

      less_corr_train_data = train_data.drop(high_corr_features, axis=1).copy()
      less_corr_test_data = test_data.drop(high_corr_features, axis=1).copy()

      less_corr_train_data.to_csv("data/less_corr_train_data.csv", index=False)
      less_corr_test_data.to_csv("data/less_corr_test_data.csv", index=False)
```

Standardizing train and test data

```
[97]: # Applying StandardData to scaled data
      # Train data
      df = less_corr_train_data.drop(['churn', 'international_plan',
      → 'voice_mail_plan'], axis=1).copy()

      StandardData = pd.DataFrame(StandardScaler().fit_transform(df), \
      columns = df.columns)

      stand_train_data = StandardData
      stand_train_data['international_plan'] = less_corr_train_data.
      → international_plan.copy()
      stand_train_data['voice_mail_plan'] = less_corr_train_data.voice_mail_plan.
      → copy()
      stand_train_data['churn'] = less_corr_train_data.churn.copy()

      # Test data
      df = less_corr_test_data.drop(['churn', 'international_plan',
      → 'voice_mail_plan'], axis=1).copy()

      StandardData = pd.DataFrame(StandardScaler().fit_transform(df), \
      columns = df.columns)

      stand_test_data = StandardData
      stand_test_data['international_plan'] = less_corr_test_data.international_plan.
      → copy()
      stand_test_data['voice_mail_plan'] = less_corr_test_data.voice_mail_plan.copy()
      stand_test_data['churn'] = less_corr_test_data.churn.copy()

      # Saving data
      stand_train_data.to_csv("data/stand_train_data.csv", index=False)
      stand_test_data.to_csv("data/stand_test_data.csv", index=False)
```

Applying MinMaxScaler to train and test data

```
[98]: # Applying MinMaxData to scaled data
# Train data
df = less_corr_train_data.drop(['churn', 'international_plan',
    ↳ 'voice_mail_plan'], axis=1).copy()

MinMaxData = pd.DataFrame(MinMaxScaler().fit_transform(df), \
                           columns = df.columns)

minmax_train_data = MinMaxData
minmax_train_data['international_plan'] = less_corr_train_data.
    ↳ international_plan.copy()
minmax_train_data['voice_mail_plan'] = less_corr_train_data.voice_mail_plan.
    ↳ copy()
minmax_train_data['churn'] = less_corr_train_data.churn.copy()

# Test data
df = less_corr_test_data.drop(['churn', 'international_plan',
    ↳ 'voice_mail_plan'], axis=1).copy()

MinMaxData = pd.DataFrame(MinMaxScaler().fit_transform(df), \
                           columns = df.columns)

minmax_test_data = MinMaxData
minmax_test_data['international_plan'] = less_corr_test_data.international_plan.
    ↳ copy()
minmax_test_data['voice_mail_plan'] = less_corr_test_data.voice_mail_plan.copy()
minmax_test_data['churn'] = less_corr_test_data.churn.copy()

# Saving data
minmax_train_data.to_csv("data/minmax_train_data.csv", index=False)
minmax_test_data.to_csv("data/minmax_test_data.csv", index=False)
```

1.4 Training models

To predict Customer Churn, I chose to use Logistic Regression to get information on whether customers are going to cancel their plan and their likelihood. ### Balancing data

```
[115]: # Defining features
Xo_train, Xo_test, yo_train, yo_test = less_corr_train_data.drop("churn",
    ↳ axis=1).copy(), \
                                     less_corr_test_data.drop("churn",
    ↳ axis=1).copy(), \
                                     less_corr_train_data.churn.copy(),
    ↳ less_corr_test_data.churn.copy()
```

1.4.1 Logistic Regression

```
[146]: # LogisticRegression algorithm
lr = LogisticRegression(C=1e5)

# training model
lr.fit(Xo_train, yo_train)

# prediction
churn_prob = lr.predict(Xo_test)
pred = churn_prob.round()

# Evaluating prediction
print (confusion_matrix(yo_test,pred))
print (classification_report(yo_test,pred))

# Save the model as a pickle in a file
# joblib.dump(lr, 'lr.pkl')
```

```
[[1418   25]
 [ 192   32]]
```

		precision	recall	f1-score	support
	0	0.88	0.98	0.93	1443
	1	0.56	0.14	0.23	224
accuracy				0.87	1667
macro avg		0.72	0.56	0.58	1667
weighted avg		0.84	0.87	0.83	1667

```
[146]: ['lr.pkl']
```

1.5 Trying to optimize model

1.5.1 Logistic regression optimization

Selecting most important features

```
[149]: # selecting most important features
k = 9

X = less_corr_train_data.drop("churn", axis=1).copy()
y = less_corr_train_data.churn.copy()

selectChi2 = SelectKBest(chi2, k=k).fit(X, y)
selectF_classif = SelectKBest(f_classif, k=k).fit(X, y)

selected = selectChi2.get_support() & selectF_classif.get_support()
```

```
features = [col for col, i in zip(X.columns, selected) if i]

print("{} most important features: ".format(len(features)), features)
```

```
7 most important features: ['international_plan', 'voice_mail_plan',
'number_vmail_messages', 'total_day_charge', 'total_eve_charge',
'total_intl_calls', 'number_customer_service_calls']
```

Training model using less variables

```
[172]: X_train, y_train = less_corr_train_data[features].copy(), less_corr_train_data.
        ↪ churn.copy()
X_test, y_test = less_corr_test_data[features].copy(), less_corr_test_data.
        ↪ churn.copy()

# LogisticRegression algorithm
lr = LogisticRegression(C=1e5)

# training model
lr.fit(X_train, y_train)

# prediction
churn_prob = lr.predict(X_test)
pred = churn_prob.round()

# Evaluating prediction
print (confusion_matrix(y_test,pred))
print (classification_report(y_test,pred))
print(accuracy_score(y_test,pred))
```

```
[[1410   33]
 [ 184   40]]

      precision    recall  f1-score   support

     0       0.88       0.98       0.93       1443
     1       0.55       0.18       0.27        224

 accuracy                   0.87       1667
 macro avg       0.72       0.58       0.60       1667
 weighted avg    0.84       0.87       0.84       1667
```

```
0.8698260347930414
```

Using standard deviation dataset

```
[171]: # Defining features
Xsd_train, Xsd_test, ysd_train, ysd_test = stand_train_data.drop("churn",
        ↪ axis=1).copy(), \
```



```

stand_test_data.drop("churn", axis=1).
→copy(), \
stand_train_data.churn.copy(),

→stand_test_data.churn.copy()

# LogisticRegression algorithm
lr = LogisticRegression(C=1e5)

# training model
lr.fit(Xsd_train, ysd_train)

# prediction
churn_prob = lr.predict(Xsd_test)
pred = churn_prob.round()

# Evaluating prediction
print (confusion_matrix(ysd_test,pred))
print (classification_report(ysd_test,pred))

print(accuracy_score(ysd_test,pred))

# Save the model as a pickle in a file
joblib.dump(lr, 'lr.pkl')

```

```

[[1407   36]
 [ 181   43]]

```

	precision	recall	f1-score	support
0	0.89	0.98	0.93	1443
1	0.54	0.19	0.28	224
accuracy			0.87	1667
macro avg	0.72	0.58	0.61	1667
weighted avg	0.84	0.87	0.84	1667

0.8698260347930414

optimization using GridSearch

```

[155]: #Results dataframe
cols = ['Case', 'LogReg']

resul = pd.DataFrame(columns=cols)
resul.set_index("Case",inplace=True)

resul.loc['Standard'] = [0]
resul.loc['GridSearch'] = [0]
resul.loc['RandomSearch'] = [0]

```

```
resul.loc['Hyperopt'] = [0]
```

[164]: *#Models creation*

```
lr = LogisticRegression(solver='liblinear')
```

```
# Standard parameters
```

```
lr.fit(Xsd_train,ysd_train.values.ravel())
```

```
resul.iloc[0, 0] = lr.score(Xsd_test,ysd_test)
```

[174]: *# GridSearch parameters*

```
from sklearn.model_selection import GridSearchCV
```

```
from sklearn.model_selection import RepeatedStratifiedKFold
```

```
#Logistic Regrresion
```

```
solvers = ['newton-cg', 'lbfgs', 'liblinear']
```

```
penalty = ['l2']
```

```
c_values = [100, 10, 1.0, 0.1, 0.01]
```

```
lr_grid = dict(solver=solvers,penalty=penalty,C=c_values)
```

```
grid = lr_grid
```

```
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3,  
                             random_state=1)
```

```
grid_search = GridSearchCV(estimator=lr,  
                           param_grid=grid, n_jobs=-1, cv=cv,  
                           scoring='accuracy',error_score=0)
```

```
grid_clf_acc = grid_search.fit(Xsd_train, ysd_train)
```

```
resul.iloc[1,0] = grid_clf_acc.score(Xsd_test,ysd_test)
```

```
# Save the model as a pickle in a file
```

```
joblib.dump(grid_search, 'grid_search.pkl')
```

[174]: ['grid_search.pkl']

optimization using RandomSearch

[165]: *from scipy.stats import randint as sp_randint*

```
from sklearn.model_selection import RandomizedSearchCV
```

```
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3,  
                             random_state=1)
```

```
n_iter_search = 3
```

```
random_search = RandomizedSearchCV(lr, param_distributions=grid,  
                                   n_iter=n_iter_search, cv=cv)
```

```
random_search.fit(Xsd_train, ysd_train)
resul.iloc[2,0] = random_search.score(Xsd_test,ysd_test)
```

```
[166]: resul.head()
```

```
[166]:          LogReg
Case
Standard      0.870426
GridSearch     0.871626
RandomSearch   0.869826
Hyperopt              0
```

1.5.2 Final model

GridSearch exibed the best result. The final model presents accuracy of 0.87. The results are shown in the table below

```
[191]: # Final model
pred = grid_search.predict(Xsd_test).round()

churn_prob = grid_search.predict_proba(Xsd_test)[: ,1]

result = pd.DataFrame({'real_churn': ysd_test, 'predicted_churn': pred,
                      'probability': churn_prob})

print (confusion_matrix(ysd_test,pred))
print (classification_report(ysd_test,pred))
print(accuracy_score(ysd_test,pred))
```

```
[[1410   33]
 [ 181   43]]

              precision    recall  f1-score   support

         0       0.89        0.98        0.93        1443
         1       0.57        0.19        0.29         224

 accuracy                   0.87        1667
 macro avg       0.73        0.58        0.61        1667
weighted avg       0.84        0.87        0.84        1667
```

```
0.871625674865027
```

```
[197]: display(result)
```

```
      real_churn  predicted_churn  probability
0              0                0      0.066634
1              0                0      0.051263
```

2	0	0	0.239797
3	0	0	0.133904
4	0	0	0.063361
...
1662	0	0	0.183550
1663	1	0	0.402079
1664	0	0	0.051855
1665	0	0	0.031570
1666	0	0	0.004603

[1667 rows x 3 columns]

```
[198]: result.to_csv("data/result.csv", index=False)
```